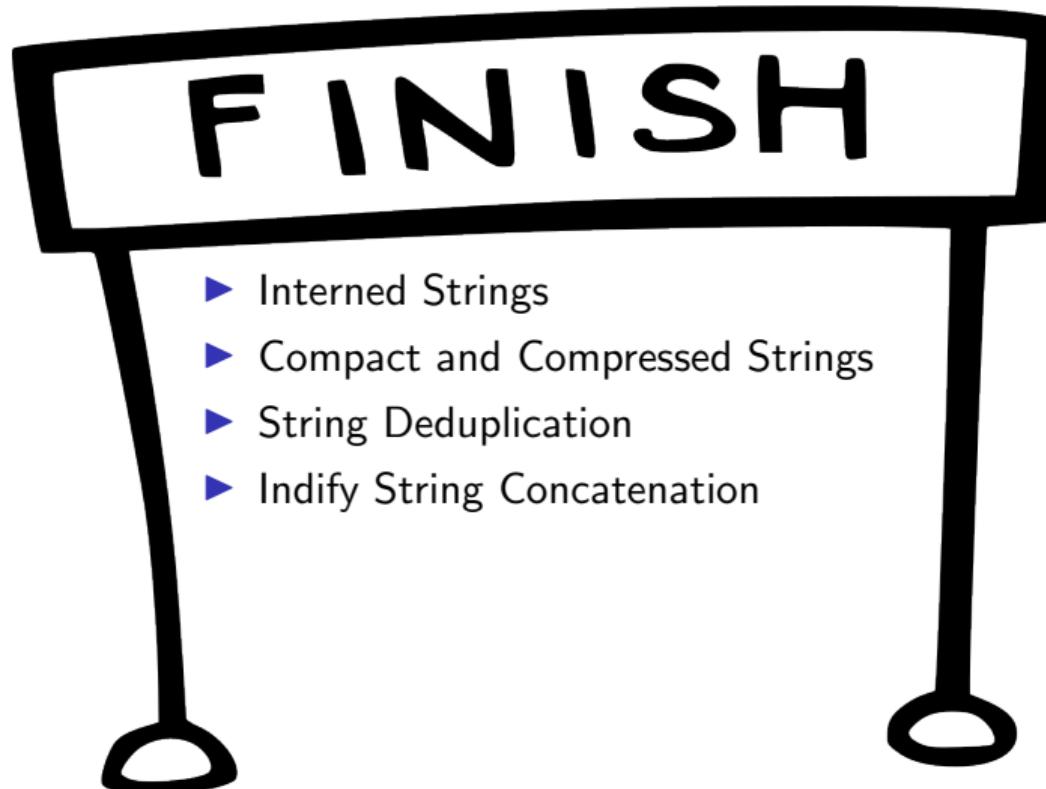


Was jeder Java-Entwickler über Strings wissen sollte

Bernd Müller
Ostfalia





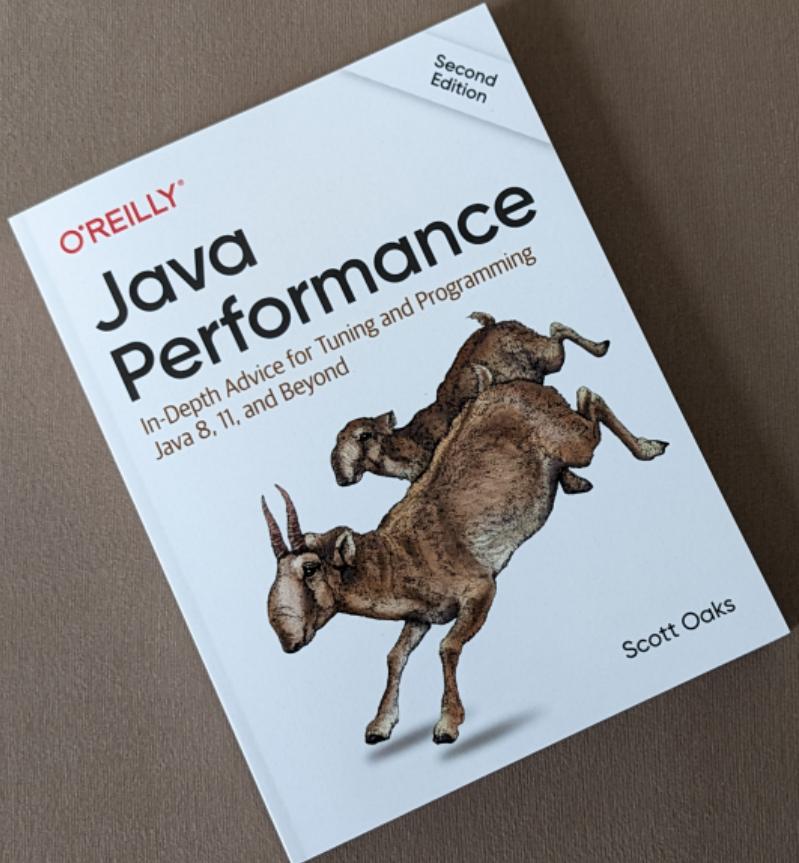
Speaker

- ▶ Prof. Computer Science (Ostfalia, HS Braunschweig/Wolfenbüttel)
- ▶ Book author (JSF, JPA, Seam, ...)



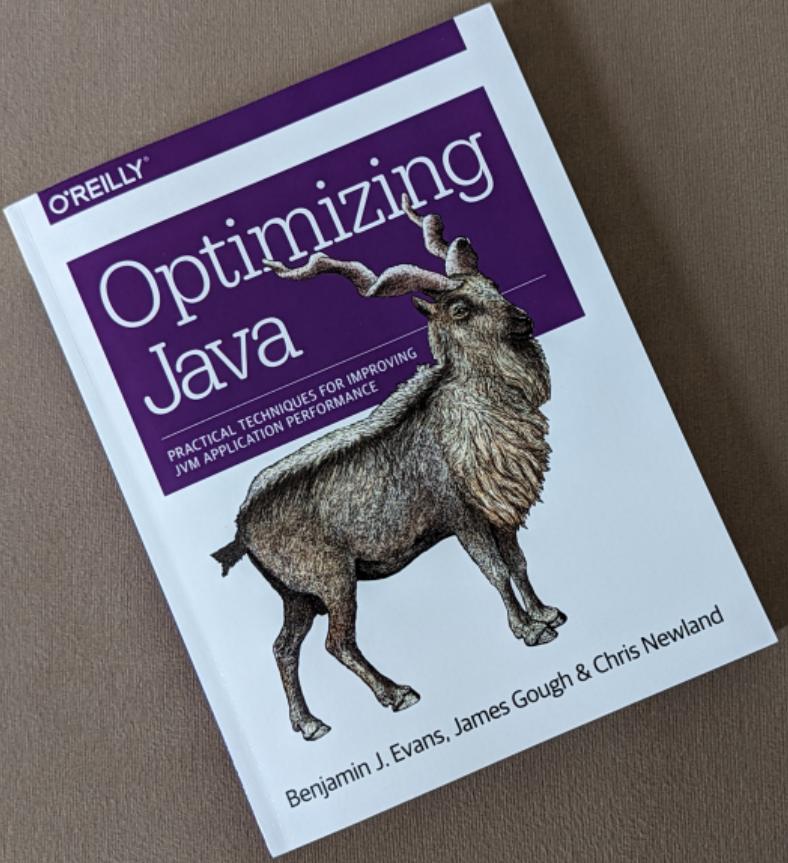
- ▶ Member EGs JSR 344 (JSF 2.2) und JSR 338 (JPA 2.1)
- ▶ Managing Director PMST GmbH
- ▶ JUG Ostfalen Co-Organizer
- ▶ bernd.mueller@ostfalia.de
- ▶ @berndmuller @berndmuller@fosstodon.org BerndMuller

Motivation



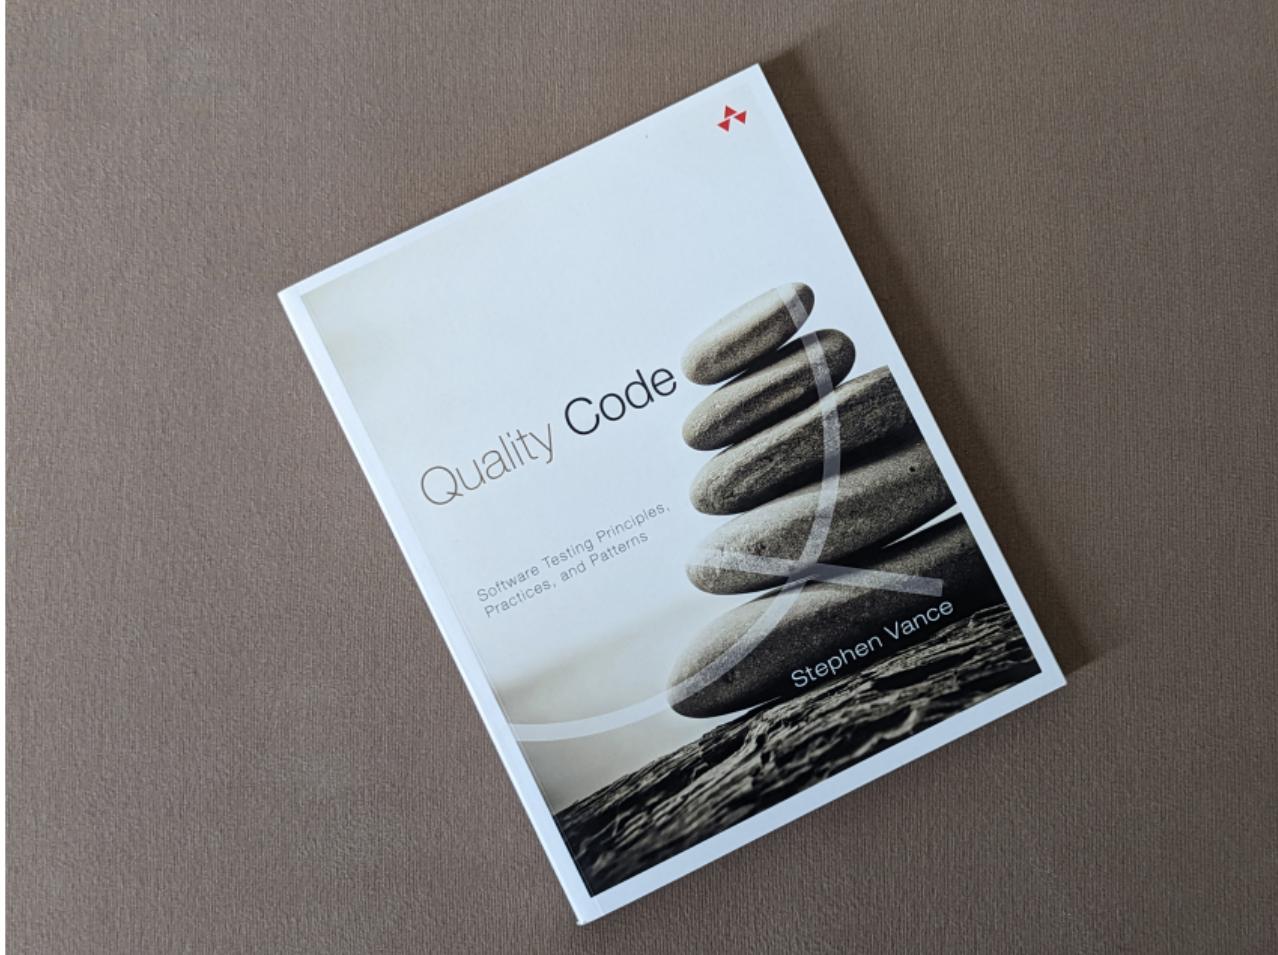
Scott Oaks, Java Performance

| | |
|--|-----|
| 12. Java SE API Tips..... | 363 |
| Strings | 363 |
| Compact Strings | 363 |
| Duplicate Strings and String Interning | 364 |
| String Concatenation | 371 |
| Buffered I/O | 374 |
| Classloading | 377 |
| Class Data Sharing | 377 |
| Random Numbers | 381 |
| Java Native Interface | 383 |
| Exceptions | 386 |



Benjamin Evens et al., Optimizing Java

| | |
|--|------------|
| 15. Java 9 and the Future..... | 377 |
| Small Performance Enhancements in Java 9 | 378 |
| Segmented Code Cache | 378 |
| Compact Strings | 378 |
| New String Concatenation | 379 |
| C2 Compiler Improvements | 381 |
| New Version of G1 | 382 |
| Java 10 and Future Versions | 382 |
| New Release Process | 383 |
| Java 10 | 383 |
| Unsafe in Java 9 and Beyond | 385 |



Stephen Vance, Quality Code

| | |
|---|----|
| Encapsulate and Override | 72 |
| Adjust Visibility | 75 |
| Verification by Injection | 77 |
| Chapter 7: String Handling | |
| Verification by Containment | 81 |
| Verification by Pattern | 83 |
| Exact Verification by Value | 85 |
| Exact Verification with Formatted Results | 88 |
| Chapter 8: Encapsulation and Override Variations | |
| Data Injection | 91 |
| Ensuring Java Conditions | 94 |

Let's start with a quiz



How many bytes are used to store a single character inside a string ?

- (a) 1 byte
- (b) 2 bytes
- (c) 4 bytes
- (d) 2 or 4 bytes
- (e) 1, 2 or 4 bytes

String Pool and String Interning

Java Language Specification

3.10.5 String Literals

...

Moreover, a string literal always refers to the same instance of class String. This is because **string literals** - or, more generally, strings that are the values of **constant expressions** (§15.29) - are "interned" so as to share unique instances, using the method `String.intern` (§12.5).

...

Java Language Specification

12.5 Creation of New Class Instances

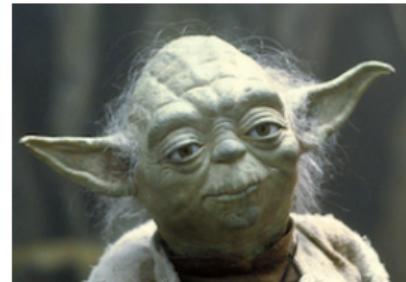
...
Loading of a class or interface that contains a string literal (§3.10.5) or a text block (§3.10.6) **may create a new String object** to denote the string represented by the string literal or text block. (**This object creation will not occur if an instance of String denoting the same sequence of Unicode code points as the string represented by the string literal or text block has previously been interned.**)

...
Execution of a string concatenation operator + (§15.18.1) that is not part of a constant expression (§15.29) always creates a new String object to represent the result. String concatenation operators may also create temporary wrapper objects for a value of a primitive type.

Summing up

- ▶ All interned strings are stored in a string memory pool
- ▶ When a class gets loaded check if string already in the pool
- ▶ If so use it, if not insert it
- ▶ Result: each string literal is a singleton inside JVM
- ▶ Also summed up in JavaDoc of `String.intern()`

Demo Time



Sleepy from slides, we are !

Performance Tip from Scott Oaks [Java Performance]

“On the topic of interning strings, what about using the `intern()` method to make the programm run faster, since interned strings can be compared via the `==` operator? That is a popular thought, though in most cases it turns out to by a myth. The `String.equals()` method is pretty fast.”

Performance Tip from Scott Oaks [Java Performance] (cont'd)

“Comparing strings via the == operator is undeniably faster, but the cost of interning the string must also be taken into consideration.”

Performance Tip from Scott Oaks [Java Performance] (cont'd)

"Like most optimizations, interning strings shouldn't be done arbitrarily, but it can be effective if there are lots of duplicate strings occupying a significant portion of the heap."

String Pool Details

- ▶ Hash table in native memory, strings itself on heap

String Pool Details

- ▶ Hash table in native memory, strings itself on heap
- ▶ Fixed size (1009 buckets < 7. Then 60,013. Since Java 11 65,536)

String Pool Details

- ▶ Hash table in native memory, strings itself on heap
- ▶ Fixed size (1009 buckets < 7. Then 60,013. Since Java 11 65,536)
- ▶ < 6 in PermGen, ≥ 7 Heap. OOME PermGen space or OOME heap space

String Pool Details

- ▶ Hash table in native memory, strings itself on heap
- ▶ Fixed size (1009 buckets < 7. Then 60,013. Since Java 11 65,536)
- ▶ < 6 in PermGen, \geq 7 Heap. OOME PermGen space or OOME heap space
- ▶ If you want to optimize:
 - ▶ `jmap -heap <process-id>`
 - ▶ `-XX:+PrintStringTableStatistics`
 - ▶ `-XX:StringTableSize=<value>`

Off topic: Mistakes can happen ...

```
/** The offset is the first index of the storage that is used. */
private final int offset;
/** The count is the number of characters in the String. */
private final int count;
```

- ▶ used for example in `substring(begin, end)`

Off topic: Mistakes can happen ...

```
/** The offset is the first index of the storage that is used. */
private final int offset;
/** The count is the number of characters in the String. */
private final int count;
```

- ▶ used for example in `substring(begin, end)`
- ▶ [JDK-4637640 : Memory leak due to `String.substring\(\)` implementation](#)
- ▶ offset and count removed with Java 7

Compact Strings (JEP 254)

Some History

- ▶ Java started 1995 as an Internet Language
- ▶ Therefore Unicode Standard, 16 bit char type
- ▶ Class String internally:

```
/** The value is used for character storage. */
private final char value[];
```

Some History

- ▶ Java started 1995 as an Internet Language
- ▶ Therefore Unicode Standard, 16 bit char type
- ▶ Class String internally:

```
/** The value is used for character storage. */
private final char value[];
```

- ▶ With Java 5 Unicode 4.0 used: sometimes requires 32 bit

Some History (cont'd)

- ▶ So called **supplementary characters** represented as pair of `char` values for characters greater than U+FFFF. Documented in class `Character`

Some History (cont'd)

- ▶ So called **supplementary characters** represented as pair of `char` values for characters greater than U+FFFF. Documented in class `Character`
- ▶ E.g. JavaDoc of `String.length()` changed from Java 5

Returns the length of this string. The length is equal to the number of 16-bit Unicode characters in the string

to Java 6

Returns the length of this string. The length is equal to the number of Unicode code units in the string

Some History (cont'd)

- ▶ So called **supplementary characters** represented as pair of `char` values for characters greater than U+FFFF. Documented in class `Character`
- ▶ E.g. JavaDoc of `String.length()` changed from Java 5

Returns the length of this string. The length is equal to the number of 16-bit Unicode characters in the string

to Java 6

Returns the length of this string. The length is equal to the number of Unicode code units in the string

- ▶ Recap: 1 byte suffices mostly (Latin-1), sometimes 2, sometimes 4 bytes

JEP 254: Compact Strings



Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities
JDK GA/EA Builds
Mailing lists
Wiki · IRC
Bylaws · Census
Legal
JEP Process
Source code
Mercurial
GitHub
Tools
Git
jtreg harness
Groups
(overview)
Adoption
Build
Client Libraries
Compatibility &
Specification
Review
Compiler
Conformance
Core Libraries
Governing Board
HotSpot
INF Tronlin & Sunnort

JEP 254: Compact Strings

| | |
|--------------------|---|
| <i>Author</i> | Brent Christian |
| <i>Owner</i> | Xueming Shen |
| <i>Type</i> | Feature |
| <i>Scope</i> | Implementation |
| <i>Status</i> | Closed / Delivered |
| <i>Release</i> | 9 |
| <i>Component</i> | core-libs/java.lang |
| <i>Discussion</i> | core dash libs dash dev at openjdk dot java dot net |
| <i>Effort</i> | L |
| <i>Duration</i> | XL |
| <i>Relates to</i> | JEP 192: String Deduplication in G1 8144691: JEP 254: Compact Strings: endiannes mismatch in Java source code and intrinsic JEP 250: Store Interned Strings in CDS Archives JEP 280: Indify String Concatenation |
| <i>Reviewed by</i> | Aleksey Shipilev, Brian Goetz, Charlie Hunt |
| <i>Endorsed by</i> | Brian Goetz |
| <i>Created</i> | 2014/08/04 21:54 |
| <i>Updated</i> | 2022/04/11 23:06 |

JEP 254: Compact Strings

OpenJDK

JEP 254: Compact Strings

Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities
JDK GA/EA Builds
Mailing lists
Wiki · IRC

Author Brent Christian
Owner Xueming Shen
Type Feature
Scope Implementation
Status Closed / Delivered

Summary

Adopt a more space-efficient internal representation for strings.

Adoption
Build
Client Libraries
Compatibility &
Specification
Review
Compiler
Conformance
Core Libraries
Governing Board
HotSpot
INF Trunkline & Sunspot

JEP 254: Mainly String Compression
Reviewed by Aleksey Shipilev, Brian Goetz, Charlie Hunt
Endorsed by Brian Goetz
Created 2014/08/04 21:54
Updated 2022/04/11 23:06

JEP 254: Compact Strings

OpenJDK

JEP 254: Compact Strings

Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities
IDK GA/EA Builds

Author Brent Christian
Owner Xueming Shen
Type Feature

Comments Implementation

Goals

Improve the space efficiency of the String class and related classes while maintaining performance in most scenarios and preserving full compatibility for all related Java and native interfaces.

Adoption
Build
Client Libraries
Compatibility & Specification Review
Compiler
Conformance
Core Libraries
Governing Board
HotSpot
INF Trunkline & Sunspot

JEP 280: Indify String Concatenation
Reviewed by Aleksey Shipilev, Brian Goetz, Charlie Hunt
Endorsed by Brian Goetz
Created 2014/08/04 21:54
Updated 2022/04/11 23:06

JEP 254: Compact Strings

[OpenJDK](#) [JEP 254: Compact Strings](#)

Motivation

The current implementation of the String class stores characters in a char array, using two bytes (sixteen bits) for each character. Data gathered from many different applications indicates that strings are a major component of heap usage and, moreover, that most String objects contain only Latin-1 characters. Such characters require only one byte of storage, hence half of the space in the internal char arrays of such String objects is going unused.

[Adoption](#)
[Build](#)
[Client Libraries](#)
[Compatibility & Specification Review](#)
[Compiler](#)
[Conformance](#)
[Core Libraries](#)
[Governing Board](#)
[HotSpot](#)
[INF Tronlin & Sunspot](#)

JEP 280: Indify String Concatenation
Reviewed by Aleksey Shipilev, Brian Goetz, Charlie Hunt
Endorsed by Brian Goetz
Created 2014/08/04 21:54
Updated 2022/04/11 23:06

With Java 9 it happened

```
private final char value[];
```

With Java 9 it happened

```
private final char value[];
```

changed to

```
private final byte[] value;
```

```
/**  
 * The identifier of the encoding used to encode the bytes in  
 * {@code value}. The supported values in this implementation are  
 * LATIN1, UTF16  
 */  
private final byte coder;
```

A Masterpiece of Software Engineering

No API changes in class String from Java 8 to 9

Usage

- ▶ Nothing to do. Default since Java 9
- ▶ VM parameter: `-XX:-CompactStrings`

Demo Time



Sleepy from slides, we are !

But things can also go wrong ...

- ▶ JDK 6 introduced *Compressed Strings*
- ▶ Similar idea but different implementation (two variants)
- ▶ Decision on JVM level: `-XX:+UseCompressedStrings`
- ▶ But ...

But things can also go wrong ...

- ▶ JDK 6 introduced *Compressed Strings*
- ▶ Similar idea but different implementation (two variants)
- ▶ Decision on JVM level: `-XX:+UseCompressedStrings`
- ▶ But ...
- ▶ Aleksey Shipilev: "*UseCompressedStrings was really the experimental feature, that was ultimately limited by design, error-prone, and hard to maintain.*"

But things can also go wrong ...

- ▶ JDK 6 introduced *Compressed Strings*
- ▶ Similar idea but different implementation (two variants)
- ▶ Decision on JVM level: `-XX:+UseCompressedStrings`
- ▶ But ...
- ▶ Aleksey Shipilev: "*UseCompressedStrings was really the experimental feature, that was ultimately limited by design, error-prone, and hard to maintain.*"
- ▶ Revoked with JDK 7

String Deduplication in G1 (JEP 192)

JEP 192: String Deduplication in G1



Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities
JDK GA/EA Builds

Mailing lists
Wiki ·IRC
Bylaws · Census
Legal

JEP Process

Source code
Mercurial
GitHub

Tools
Git
Jtreg harness

Groups
(overview)
Adoption
Build
Client Libraries

JEP 192: String Deduplication in G1

| | |
|--------------------|--|
| <i>Owner</i> | Per Liden |
| <i>Type</i> | Feature |
| <i>Scope</i> | Implementation |
| <i>Status</i> | Closed / Delivered |
| <i>Release</i> | 8u20 |
| <i>Component</i> | hotspot/gc |
| <i>Discussion</i> | hotspot dash gc dash dev at openjdk dot java dot net |
| <i>Effort</i> | M |
| <i>Duration</i> | L |
| <i>Relates to</i> | JEP 254: Compact Strings |
| <i>Reviewed by</i> | Bengt Rutisson, John Coomes, Jon Masamitsu |
| <i>Endorsed by</i> | Mikael Vidstedt |
| <i>Created</i> | 2013/11/22 20:00 |
| <i>Updated</i> | 2017/06/07 22:25 |
| <i>Issue</i> | 8046182 |

JEP 192: String Deduplication in G1



Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities

JEP 192: String Deduplication in G1

Owner Per Liden
Type Feature
Scope Implementation

Summary

Reduce the Java heap live-data set by enhancing the G1 garbage collector so that duplicate instances of String are automatically and continuously deduplicated.

SEARCH
Tools
Git
Jtreg harness
Groups
(overview)
Adoption
Build
Client Libraries

RELATED TO JEP 427: Compact Strings

Reviewed by Bengt Rutisson, John Coomes, Jon Masamitsu
Endorsed by Mikael Vidstedt
Created 2013/11/22 20:00
Updated 2017/06/07 22:25
Issue [8046182](#)

JEP 192: String Deduplication in G1

OpenIDK

IEP 192: String Deduplication in G1

Motivation

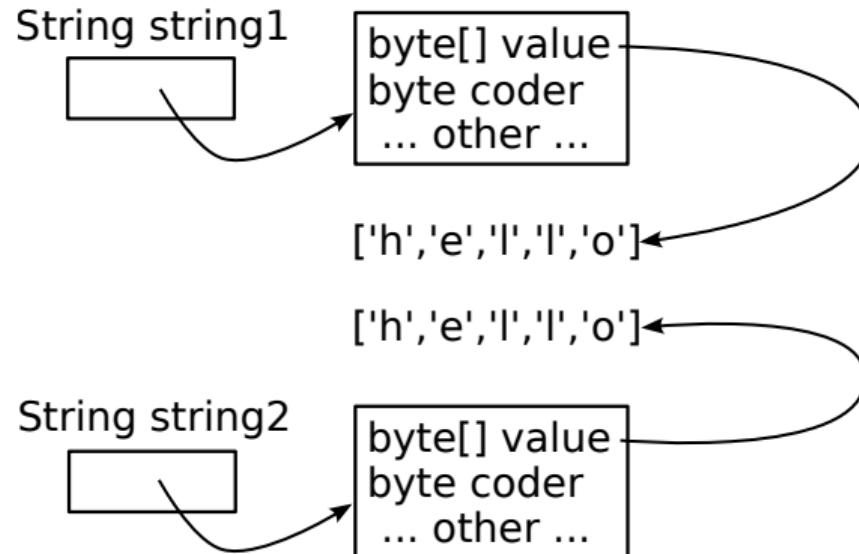
Many large-scale Java applications are currently bottlenecked on memory.

Measurements have shown that roughly 25% of the Java heap live data set in these types of applications is consumed by String objects. Further, roughly half of those String objects are duplicates, where duplicates means `string1.equals(string2)` is true. Having duplicate String objects on the heap is, essentially, just a waste of memory. This project will implement automatic and continuous String deduplication in the G1 garbage collector to avoid wasting memory and reduce the memory footprint.

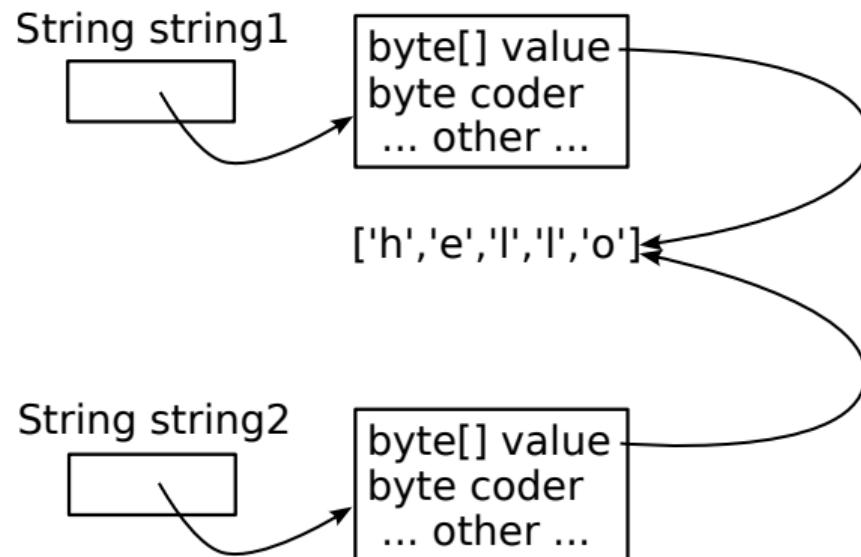
Groups
(overview)
Adoption
Build
Client Libraries

Created 2013/11/22 20:00
Updated 2017/06/07 22:25
Issue 8046182

How does it work ?



After String Deduplication



Usage

- ▶ Only works with G1 garbage collector
- ▶ VM parameter: `-XX:+UseStringDeduplication`
- ▶ Available since Java 8 update 20
- ▶ `-XX:+PrintStringDeduplicationStatistics` removed with Java 9
- ▶ Use `-Xlog:stringdedup*=debug` instead

Demo Time



Sleepy from slides, we are !

Case Study

- ▶ Article G1: from garbage collector to waste management consultant
- ▶ Case study with Eclipse IDE
- ▶ Result: Decrease heap usage by about 10%
- ▶ Heinz Kabutz reports about 25%

Indify String Concatenation (JEP 280)



Main Page
Community portal
Preferences
Requested entries
Recent changes
Random entry
Help
Glossary
Donations
Contact us

Tools

[What links here](#)

Entry [Discussion](#) [Citations](#)

Not logged in [Talk](#) [Contributions](#) [Preferences](#) [Create account](#) [Log in](#)

[Read](#) [Edit](#) [History](#)

Search Wiktionary



indify

English [\[edit\]](#)

Verb [\[edit\]](#)

indify (*third-person singular simple present **indifies**, present participle **indifying**, simple past and past participle **indified***)

1. (*programming, JAVA*) Change a particular functionality to use invokedynamic calls.

Anagrams [\[edit\]](#)

- nidify

Categories: English lemmas | English verbs | en:Programming

JEP 280: Indify String Concatenation

OpenJDK

Installing
Contributing
Sponsoring
Developers' Guide
Vulnerabilities
JDK GA/EA Builds
Mailing lists
Wiki · IRC

Bylaws · Census
Legal
JEP Process

Source code
Mercurial
GitHub
Tools
Git
Jtreg harness

Groups
(overview)
Adoption
Build
Client Libraries
Compatibility &
Specification
Review
Compiler
Conformance
Core Libraries

JEP 280: Indify String Concatenation

| | |
|--------------------|--|
| <i>Owner</i> | Aleksey Shipilev |
| <i>Type</i> | Feature |
| <i>Scope</i> | SE |
| <i>Status</i> | Closed / Delivered |
| <i>Release</i> | 9 |
| <i>Component</i> | tools/javac |
| <i>Discussion</i> | core dash libs dash dev at openjdk dot java dot net, compiler dash dev at openjdk dot java dot net, hotspot dash dev at openjdk dot java dot net |
| <i>Effort</i> | M |
| <i>Duration</i> | M |
| <i>Relates to</i> | JEP 254: Compact Strings |
| <i>Reviewed by</i> | Michael Haupt, Paul Sandoz |
| <i>Endorsed by</i> | Brian Goetz |
| <i>Created</i> | 2015/06/04 08:13 |
| <i>Updated</i> | 2022/04/28 05:04 |
| <i>Issue</i> | 8085796 |

JEP 280: Indify String Concatenation

[OpenJDK](#)

JEP 280: Indify String Concatenation

Installing
Contributing
Sponsoring
Developers' Guide

Owner Aleksey Shipilev
Type Feature
--

Summary

Change the static String-concatenation bytecode sequence generated by javac to use invokedynamic calls to JDK library functions. This will enable future optimizations of String concatenation without requiring further changes to the bytecode emitted by javac.

Groups
(overview)
Adoption
Build
Client Libraries
Compatibility &
Specification
Review
Compiler
Conformance
Core Libraries

REVIEWED [MICHAEL HEGE, ERIC SARKOZ](#)
by
Endorsed Brian Goetz
by
Created 2015/06/04 08:13
Updated 2022/04/28 05:04
Issue [8085796](#)

JEP 280: Indify String Concatenation

[OpenJDK](#)

JEP 280: Indify String Concatenation

[Installing](#)
[Contributing](#)
[Sponsoring](#)
[Developers' Guide](#)

Owner Aleksey Shipilev
Type Feature
--
--

Summary

Change the static String-concatenation bytecode sequence generated by javac to use ~~invokedynamic calls to JDK library functions~~. This will enable future optimizations of String concatenation without requiring further changes to the bytecode emitted by javac.

[Groups](#)
(overview)
[Adoption](#)
[Build](#)
[Client Libraries](#)
[Compatibility & Specification Review](#)
[Compiler](#)
[Conformance](#)
[Core Libraries](#)

REVIEWED [MICHAEL HEGPE, TANU JAIN](#)
by
Endorsed Brian Goetz
by
Created 2015/06/04 08:13
Updated 2022/04/28 05:04
Issue [8085796](#)

What they did

- ▶ Replace `StringBuilder#append()` chains with
 - ▶ invokedynamic
 - ▶ bootstrap methods
 - ▶ Class `java.lang.invoke.StringConcatFactory`
- ▶ Please verify yourself with `javap -l -s -verbose <class>`

SUMMARY: NESTED | FIELD | CONSTR | METHOD

DETAIL: FIELD | CONSTR | METHOD

SEARCH Search X**Module** java.base**Package** java.lang.invoke

Class StringConcatFactory

java.lang.Object
 java.lang.invoke.StringConcatFactory

```
public final class StringConcatFactory
extends Object
```

Methods to facilitate the creation of String concatenation methods, that can be used to efficiently concatenate a known number of arguments of known types, possibly after type adaptation and partial evaluation of arguments. These methods are typically used as *bootstrap methods* for invokedynamic call sites, to support the *string concatenation* feature of the Java Programming Language.

Indirect access to the behavior specified by the provided MethodHandle proceeds in order through two phases:

1. *Linkage* occurs when the methods in this class are invoked. They take as arguments a method type describing the concatenated arguments count and types, and optionally the String *recipe*, plus the constants that participate in the String concatenation. The details on accepted recipe shapes are described further below. Linkage may involve dynamically loading a new class that implements the expected concatenation behavior. The CallSite holds the MethodHandle pointing to the exact concatenation method. The concatenation methods may be shared among different CallSites, e.g. if linkage methods produce them as pure functions.
2. *Invocation* occurs when a generated concatenation method is invoked with the exact dynamic arguments. This may occur many times for a single concatenation method. The method referenced by the behavior MethodHandle is invoked with the

It will always go on ...



Ismael Juma
@ijuma

...

A significant `String.hashCode()` performance improvement from [@cl4es](#) has been merged - a micro-benchmarks involving a string with 10k characters shows a 6x(!) improvement. The improvement is not as extreme (but still impressive) for smaller strings.

[Tweet](#) [Übersetzen](#)

openjdk/jdk

#10847 **8282664: Unroll by hand StringUTF16 and StringLatin1...**



61 comments 39 reviews 33 files +1053 -87

cl4es • October 25, 2022 ~ 76 commits



github.com

8282664: Unroll by hand StringUTF16 and StringLatin1 polynomial hash loops ...
Continuing the work initiated by @luhenry to unroll and then intrisify polynomial hash loops. I've rewired the library changes to route via a single ...

294

String.format() 3x faster in Java 17

Author: Dr Heinz M. Kabutz | Date: 2021-10-29 | Java Version: 17 | Category: [Performance](#)

Abstract: One of the most convenient ways of constructing complex Strings is with `String.format()`. It used to be excessively slow, but in Java 17 is about 3x faster. In this newsletter we discover what the difference is and where it will help you. Also when you should use `format()` instead of the plain String addition with `+`.

Welcome to the 294th edition of **The Java(tm) Specialists' Newsletter**. We had a lovely run in the rain today, followed by a dip in the sea, clocking in at 21.6 degrees celsius. That is bathwater for someone from Bantry Bay! I remember the water in Cape Town being so cold that our breath misted as my brother and I contemplated how crazy we were to spearfish in single-digit water temperatures - and that was in summer.

javaspecialists.teachable.com: Please visit our new [self-study course catalog](#) to see how you can upskill your Java knowledge.

String.format() 3x faster in Java 17

A few years ago, my friend Dmitry Vyazelenko and I submitted a talk to JavaOne, where we spoke for about an hour about the humble `java.lang.String`. We have since spoken about this fundamental class at Devoxx, Geecon, Geekout, JAX, Voxxed Days, GOTO, and various JUGs around the world. Who would have thought that we could easily fill an hour with a talk about `java.lang.String`?

I would usually start the talk by showing a quiz. Which method is the fastest at appending Strings?

```
public class StringAppendingQuiz {  
    public String appendPlain(String question,  
                            String answer1,  
                            ...)
```

 **Subscribe Now**

Get: every new article, exclusive invites to live webinars, our top-10 articles ever, access to expert tutorials & more!

First Name *

Last Name

Company

Email *

Country

Java Skill Level *

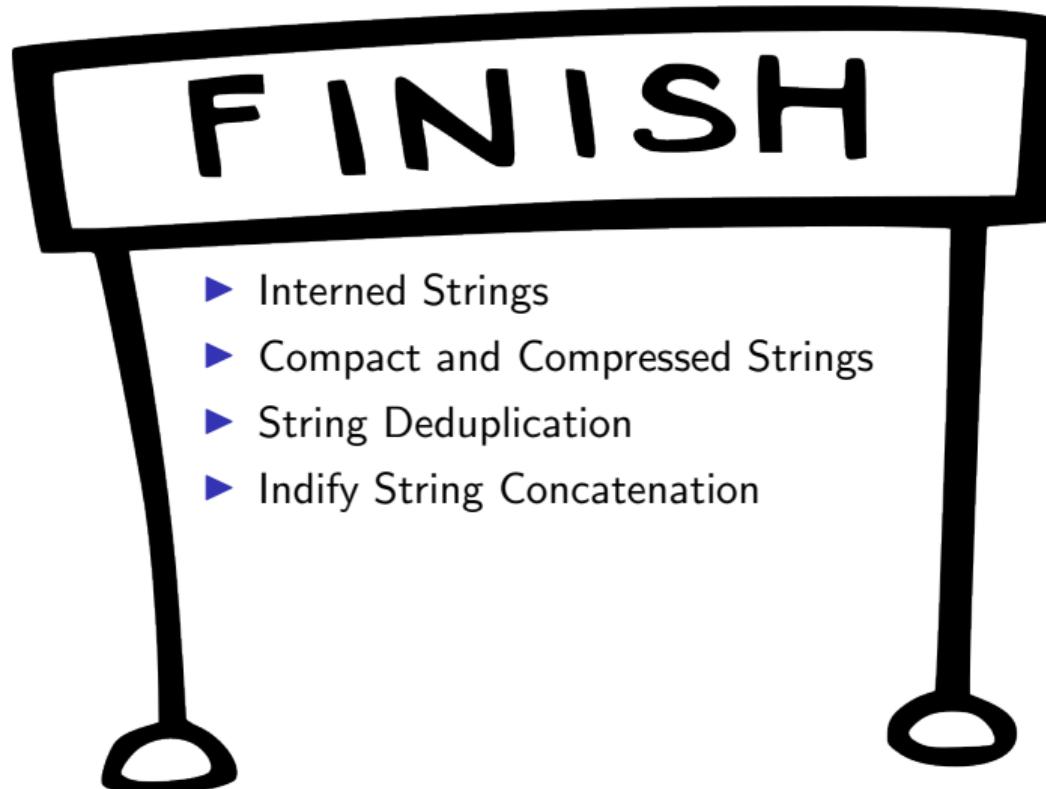
Agree to JavaSpecialists Privacy Policy *

Subscribe

[Read our Privacy Policy](#)

About the Author







Slides and Code

<https://github.com/BerndMuller/strings-jfs-2023>

